

# GPU-Quicksort: from OpenCL to Data Parallel C++

By Robert Ioffe

November 5, 2019

## Contents

Introduction .....	1
Acknowledgements.....	2
What is GPU-Quicksort?.....	2
What is OpenCL? .....	2
What is Data Parallel C++? .....	2
The Starting Point: Windows Apps from 2014.....	2
Data Parallel C++/OpenCL Interoperability: Platform Initialization.....	3
Data Parallel C++: How to Select an Intel GPU?.....	4
Data Parallel C++: How to Set Kernel Arguments and Launch Kernels? .....	5
Data Parallel C++: How to Create Buffers and Set the Access Mode?.....	6
Data Parallel C++: How to Query Platform and Device Properties? .....	6
Porting OpenCL Kernels to Data Parallel C++: Part 1 – gqsort_kernel.....	7
Porting OpenCL Kernels to Data Parallel C++: Part 2 – lqsort_kernel.....	10
The Power of Data Parallel C++: Templates! And Their Caveats .....	11
Back to Windows! And RHEL .....	12
Conclusion.....	13
References .....	13

## Introduction

Intel is about to introduce to the world Data Parallel C++: a heterogeneous portable programming language based on the Khronos SYCL standard. The promise of Data Parallel C++ is a single source programming language that can target an array of platforms that Intel has to offer: CPUs, integrated and discrete GPUs, FPGAs and other accelerators. We decided to kick the tires of this new programming language and tools that Intel offers for it and port a non-trivial OpenCL application, GPU-Quicksort, to Data Parallel C++ and document the experience. We also set the goal of exceeding the capabilities of the initial application: OpenCL C makes it very hard to write generic algorithms, and it becomes clear that it is a serious shortcoming when one tries to implement algorithms such as sorting, which you may want to do for a variety of different data types. The original GPU-Quicksort for OpenCL was written to only

sort unsigned integers. We demonstrate how to use templates with Data Parallel C++ and show how to implement GPU-Quicksort for floats and doubles as well. As a final step we try Data Parallel C++ on Windows and SLES to prove its portability.

## Acknowledgements

### What is GPU-Quicksort?

GPU-Quicksort is a high-performance sorting algorithm specifically designed for the highly parallel multicore graphics processors. It was invented in 2009 by Daniel Cederman and Phillippas Tsigas, a student/professor pair from the Chalmers University of Technology in Sweden. Originally implemented in CUDA, it was reimplemented in OpenCL 1.2 and OpenCL 2.0 in 2014 by the author of this article to demonstrate high performance on the Intel® Integrated Processor Graphics and showcase nested parallelism and work-group scan functions enabled by the OpenCL 2.0 standard and fully implemented in Intel OpenCL drivers. In this article we demonstrate how to port OpenCL 1.2 implementation of the GPU-Quicksort to Data Parallel C++ and how to make the implementation more generic and sort not only unsigned integers, but floats and doubles too.

### What is OpenCL?

We start with the OpenCL 1.2 implementation. OpenCL is a Khronos standard for programming heterogeneous systems and is fully supported by Intel on a variety of the operating systems and platforms. OpenCL consists of the runtime, the host API and the device C-based programming language OpenCL C. Here lies both its power and its limitations: the power is the ability to write high performance portable heterogeneous applications, and the main limitations are the necessity to write and debug two separate codes: the host side and the device side, and the lack of templates and other C++ features that modern programmers got accustomed to, which make writing generic libraries in OpenCL extremely hard.

### What is Data Parallel C++?

Data Parallel C++ (DPC++) is an Intel implementation of Khronos SYCL, the standard designed to address the two main limitations of OpenCL outlined above. DPC++ provides both a single source programming model, which enables to have a single code base for host and device programming and enables the full use of C++ templates and template metaprogramming on the device without the loss of portability and a very minimal impact to performance. DPC++ allows a programmer to target CPUs, GPUs and FPGAs while permitting accelerator-specific tuning, so it is a definite improvement on OpenCL. It is also supported by the Intel tools that developers come to love, such as Intel® VTune™ Profiler and Intel® Advisor as well as GDB. We want to take full advantage of DPC++, especially the its template features.

### The Starting Point: Windows Apps from 2014

We start our journey with GPU-Quicksort for OpenCL 1.2 implementation as provided in the article “GPU-Quicksort in OpenCL 2.0: Nested Parallelism and Work-Group Scan Functions” by the author. The original application was written for Windows, so we port it to Ubuntu 18.04 Linux by adding the cross-platform code to measure time and use `aligned_alloc/free` for aligned memory allocation/deallocation

as opposed to `_aligned_malloc/_aligned_free` on Windows. Now, here is a brief overview of GPU-Quicksort application architecture. It consists of two kernels: `gqsort_kernel` and `lqsort_kernel`, both written in OpenCL 1.2, that are glued together by a dispatcher code, which iteratively calls `gqsort_kernel` until the input is split into small enough chunks, which can be fully sorted by `lqsort_kernel`. The application allows the user to select the number of times to run sort for measurement purposes, the vendor and device to run the kernels on, the size of the input and whether to show device details that the code is running on. The application follows a typical OpenCL architecture of supporting utilities for initializing OpenCL platforms and devices and building code for them, a separate file with the OpenCL kernels and their supporting functions and the main application that accepts user arguments, initializes platform and device, builds the kernels, properly allocates memory and creates buffers from it, binds them to the kernel arguments and then launches the dispatcher function.

## Data Parallel C++/OpenCL Interoperability: Platform Initialization

The first thing that you need to do to try DPC++ is to install Intel® oneAPI Base Toolkit. DPC++ compiler `dpcpp` is part of that toolkit. We start out port to DPC++ by including `CL/sycl.hpp` header and to spare us the verbosity of DPC++, using `namespace cl::sycl` clause:

```
#include <CL/sycl.hpp>
...
using namespace cl::sycl;
```

Now, instead of initializing a platform, a device and a context and a queue the OpenCL way, we need to do it the DPC++ way:

```
device d(default_selector);
if (d.is_host()) {
    // This platform can't pass this test, it has no OpenCL devices
    return;
}

queue queue(default_selector, [] (cl::sycl::exception_list l) {
    for (auto ep : l) {
        try {
            std::rethrow_exception(ep);
        } catch (cl::sycl::exception e) {
            std::cout << e.what() << std::endl;
        }
    }
});
```

We also need to retrieve the underlying OpenCL context, device and queue, since the rest of the application is still very much OpenCL based:

```
/* Retrieve the underlying cl_context of the context associated with the
 * queue. */
pOCL->contextHdl = queue.get_context().get();

/* Retrieve the underlying cl_device_id of the device associated with the
```

```

    * queue. */
pOCL->deviceID = queue.get_device().get();

/* Retrieve the underlying cl_command_queue of the queue. */
pOCL->cmdQHdl = queue.get();

```

And that's is pretty much it for our first iteration: now configure and compile it with dpcpp compiler and you are ready to run it.

## Data Parallel C++: How to Select an Intel GPU?

The only shortcoming of the iteration above is that it is always selects the default device, which may or may not be an Intel GPU. If we want to be able to select Intel GPU, we need to write a custom device selector:

```

/* Classes can inherit from the device_selector class to allow users
 * to dictate the criteria for choosing a device from those that might be
 * present on a system. This example looks for a device with Intel in its name
 * and prefers GPUs. */
class intel_gpu_selector : public device_selector {
public:
    intel_gpu_selector() : device_selector() {}

    /* The selection is performed via the () operator in the base
     * selector class. This method will be called once per device in each
     * platform. Note that all platforms are evaluated whenever there is
     * a device selection. */
    int operator()(const device& device) const override {
        /* We only give a valid score to devices that support SPIR. */
        //if (device.has_extension(cl::sycl::string_class("cl_khr_spir"))) {
        if (device.get_info<info::device::name>().find("Intel") != std::string::npos) {
            if (device.get_info<info::device::device_type>() ==
                info::device_type::gpu) {
                return 50;
            }
        }
        /* Devices with a negative score will never be chosen. */
        return -1;
    }
};

```

We can now use our intel\_gpu\_selector to select an Intel GPU when the user asks for it:

```

auto get_queue = [&pDeviceStr, &pVendorStr]() {
    device_selector* pds;
    if (pVendorStr == std::string("intel")) {
        if (pDeviceStr == std::string("gpu")) {
            static intel_gpu_selector selector;
            pds = &selector;
        } else if (pDeviceStr == std::string("cpu")) {
            static cpu_selector selector;
            pds = &selector;
        }
    }
}

```

```

} else {
    static default_selector selector;
    pds = &selector;
}

device d(*pds);

queue queue(*pds, [](cl::sycl::exception_list l) {
    for (auto ep : l) {
        try {
            std::rethrow_exception(ep);
        } catch (cl::sycl::exception e) {
            std::cout << e.what() << std::endl;
        }
    }
});
return queue;
};

auto queue = get_queue();

```

## Data Parallel C++: How to Set Kernel Arguments and Launch Kernels?

The third iteration of our code is to use DPC++ to set kernel arguments and launch kernels. Note that the program is still built the old OpenCL way and the kernels are obtained from it the old way too. We use `cl::sycl::kernel` objects to wrap original OpenCL kernels, e.g.

```
cl::sycl::kernel sycl_gqsort_kernel(gqsort_kernel, pOCL->queue.get_context());
```

Then we replace a number of `clSetKernelArg` methods with `set_arg` DPC++ methods and `clEnqueueNDRange` calls with `parallel_for` calls (the example below shows `gqsort_kernel`, but `lqsort_kernel` upgrade is very similar):

```

pOCL->queue.submit([&](handler& cgh) {
    /* Normally, SYCL sets kernel arguments for the user. However, when
     * using the interoperability features, it is unable to do this and
     * the user must set the arguments manually. */
    cgh.set_arg(0, db);
    cgh.set_arg(1, دنب);
    cgh.set_arg(2, blocksb);
    cgh.set_arg(3, parentsb);
    cgh.set_arg(4, newsb);

    cgh.parallel_for(nd_range<1>(range<1>(GQSORT_LOCAL_WORKGROUP_SIZE * (blocks.size())),
                                     range<1>(GQSORT_LOCAL_WORKGROUP_SIZE)),
                     sycl_gqsort_kernel);
});
pOCL->queue.wait_and_throw();

```

We can use a less verbose style and set all the arguments of the kernel with one `set_args` call:

```
cgh.set_args(db, دنب, blocksb, parentsb, newsb);
```

We can also use a less verbose version of the `parallel_for`:

```
cgh.parallel_for(nd_range<>(GQSORT_LOCAL_WORKGROUP_SIZE * blocks.size(),
                           GQSORT_LOCAL_WORKGROUP_SIZE),
```

## Data Parallel C++: How to Create Buffers and Set the Access Mode?

Now we are ready to convert our OpenCL buffers to DPC++ buffers (the first two buffers are wrapping the memory that was align allocated and passed into the function by reference where the other three buffers are created from an STL vector). Note the use of template keyword in front of the `get_access` member function for buffers that we passed by reference. Also note different access modes for various buffers depending on whether we need read or write access or both. We don't directly pass buffers as kernel arguments: we pass the accessors to them!

```
    buffer<T> d_buffer(d, range<>(size), {property::buffer::use_host_ptr()});
    buffer<T> dn_buffer(dn, range<>(size), {property::buffer::use_host_ptr()});
    ...
    // Create buffer objects for memory.
    buffer<block_record> blocks_buffer(blocks.data(), blocks.size(),
{property::buffer::use_host_ptr()});
    buffer<parent_record> parents_buffer(parents.data(), parents.size(),
{property::buffer::use_host_ptr()});
    buffer<work_record> news_buffer(news.data(), news.size(),
{property::buffer::use_host_ptr()});

    pOCL->queue.submit([&](handler& cgh) {
        auto db = d_buffer.template get_access<access::mode::discard_read_write>(cgh);
        auto dnb = dn_buffer.template get_access<access::mode::discard_read_write>(cgh);
        auto blocksb = blocks_buffer.get_access<access::mode::read>(cgh);
        auto parentsb = parents_buffer.get_access<access::mode::read>(cgh);
        auto newsb = news_buffer.get_access<access::mode::write>(cgh);
        /* Normally, SYCL sets kernel arguments for the user. However, when
         * using the interoperability features, it is unable to do this and
         * the user must set the arguments manually. */
        cgh.set_args(db, dnb, blocksb, parentsb, newsb);

        cgh.parallel_for(nd_range<>(GQSORT_LOCAL_WORKGROUP_SIZE * blocks.size(),
                                   GQSORT_LOCAL_WORKGROUP_SIZE),
                        sycl_gqsort_kernel);
    });
    pOCL->queue.wait_and_throw();
```

## Data Parallel C++: How to Query Platform and Device Properties?

Now, in good old OpenCL days, we used the methods such as `clGetPlatformInfo` and `clGetDeviceInfo` to query various platform and device properties. Now we can use `get_info<>` methods to query the same information, e.g.

```
auto max_work_item_dimensions =
q.get_device().get_info<info::device::max_work_item_dimensions>();
```

```
std::cout << "CL_DEVICE_MAX_WORK_ITEM_DIMENSIONS: " << max_work_item_dimensions << std::endl;
```

or query properties with a more complex structure:

```
auto max_work_item_sizes = q.get_device().get_info<info::device::max_work_item_sizes>();
printf ("CL_DEVICE_MAX_WORK_ITEM_SIZES : (%5zu, %5zu, %5zu)\n",
        max_work_item_sizes[0], max_work_item_sizes[1], max_work_item_sizes[2]);
```

## Porting OpenCL Kernels to Data Parallel C++: Part 1 – gqsort\_kernel

So far we initialized the platform and the device, created the buffers, their accessors and bound them to the kernels and launched those kernel on the device in a DPC++ way, but we still create the kernels in an old-fashioned OpenCL way: we use OpenCL C and `clBuildProgram/clCreateKernel` APIs to build the program and to create our kernels. The OpenCL C kernels are stored in a separate file that is loaded into the program at runtime before being built. We are about to change that starting with the `gqsort_kernel`, arguably the simpler of the two kernels.

DPC++ way of creating kernels is via lambdas or functors. The use of lambdas for kernel creation is typically reserved for smaller kernels. When you have a more complex kernel that uses supporting functions it is a good idea to create a functor class. We are going to create a `gqsort_kernel_class` functor and make it templated right from the start to be able to handle sorting more than one datatype in the future.

```
template <class T>
class gqsort_kernel_class {
...
};
```

A typical functor class will have a `void operator()` that will take as a parameter an iteration id (in our case, a 1-dimensional `nd_item<1> id`). The body of the kernel will reside in the `void operator()`. The functor will also have a constructor that will take global and local accessors: the equivalent of global and local memory pointers for an OpenCL kernel. The typical DPC++ functor will have a preamble with `using` clauses defining various global and local accessor types. In the case of `gqsort_kernel`, it will look like this:

```
using blocks_read_accessor = accessor<block_record<T>, 1, access::mode::read,
access::target::global_buffer>;
using parents_read_write_accessor = accessor<parent_record, 1, access::mode::read_write,
access::target::global_buffer>;
using news_write_accessor = accessor<work_record<T>, 1, access::mode::write,
access::target::global_buffer>;
using discard_read_write_accessor =
    accessor<T, 1, access::mode::discard_read_write, access::target::global_buffer>;
using local_read_write_accessor = accessor<uint, 1, access::mode::read_write,
access::target::local>;
```

The private section of the functor will contain all the global and local accessors used within the body of the `void operator()`. In our case, it will look like this, with the first five accessors to global buffers and the rest to the local buffers:

```
private:
```

```
discard_read_write_accessor d, dn;
blocks_read_accessor blocks;
parents_read_write_accessor parents;
news_write_accessor news;
local_read_write_accessor lt, gt, ltsum, gtsum, lbeg, gbeg;
```

gqsort\_kernel is a relatively complex kernel that uses supporting structs and two supporting functions: plus\_prescan and median, which in turn uses specialized OpenCL functions, extensively uses local memory arrays and variables, local and global barriers and atomics. All these elements have to be translated into DPC++. Let's start with the functions (we omit structs, since they are trivially templated). Let's start with the plus\_prescan function that is used to calculate scan sums. It is a relatively simple function, so the only change that we will make to bring it to DPC++ is to make it a template function in preparation of making our sort a generic algorithm:

```
template <class T>
void plus_prescan(T *a, T *b) {
    T av = *a;
    T bv = *b;
    *a = bv;
    *b = bv + av;
}
```

We will deal with the median function next: not only do we make it a template function, but we also need to replace OpenCL C select functions with DPC++ cl::sycl::select functions. We will also rename it median\_select to differentiate it from a similar host function:

```
template <class T>
T median_select(T x1, T x2, T x3) {
    if (x1 < x2) {
        if (x2 < x3) {
            return x2;
        } else {
            return cl::sycl::select(x1, x3, (uint)(x1 < x3));
        }
    } else { // x1 >= x2
        if (x1 < x3) {
            return x1;
        } else { // x1 >= x3
            return cl::sycl::select(x2, x3, (uint)(x2 < x3));
        }
    }
}
```

Now we must deal with the local memory variables and arrays. While in OpenCL C it is possible both to create local memory variables and arrays inside the body of the kernel and pass them as kernel parameters, in DPC++ when using functors we pass local buffer accessors when constructing the functor. In our case, all local memory variable and arrays will hold unsigned integers, so we will create a special local\_read\_write\_accessor type:



```
using local_read_write_accessor = accessor<uint, 1, access::mode::read_write,
access::target::local>;
```

We then declare all the local memory variables as follows:

```
local_read_write_accessor
    lt(range<>(GQSORT_LOCAL_WORKGROUP_SIZE+1), cgh),
gt(range<>(GQSORT_LOCAL_WORKGROUP_SIZE+1), cgh),
    ltsum(range<>(1), cgh), gtsum(range<>(1), cgh), lbeg(range<>(1), cgh), gbeg(range<>(1),
cgh);
```

We then pass them as parameters to our functor constructor along with global buffer accessors, and the resulting object is passed to the `parallel_for`:

```
auto gqsort = gqsort_kernel_class<T>(db, دنب, blocksb, parentsb, newsb, lt, gt, ltsum,
gtsum, lbeg, gbeg);

cgh.parallel_for(
    nd_range<>(GQSORT_LOCAL_WORKGROUP_SIZE * blocks.size(),
              GQSORT_LOCAL_WORKGROUP_SIZE),
    gqsort);
```

Here DPC++ notably lacks in simplicity as compared to OpenCL C. Next, we deal with `get_group_id` and `get_local_id` functions, which in DPC++ become:

```
const size_t blockid = id.get_group(0);
const size_t localid = id.get_local_id(0);
```

Local barriers go from

```
barrier(CLK_LOCAL_MEM_FENCE);
```

to

```
id.barrier(access::fence_space::local_space);
```

While global and local barriers go from

```
barrier(CLK_LOCAL_MEM_FENCE | CLK_GLOBAL_MEM_FENCE);
```

to

```
id.barrier(access::fence_space::global_and_local);
```

On to atomics! Here again DPC++ is not as elegant as OpenCL C, so a relatively concise:

```
// get shared variables
psstart = &parent->sstart;
psend = &parent->send;
```

```

// Atomic increment allocates memory to write to.
lbegin = atomic_add(psstart, ltsum);

// Atomic is necessary since multiple blocks access this
gbegin = atomic_sub(psend, gtsum) - gtsum;

```

becomes fairly unwieldy:

```

cl::sycl::atomic<uint> psstart_a(multi_ptr<uint,
access::address_space::global_space>(&pparent.sstart));
cl::sycl::atomic<uint> psend_a(multi_ptr<uint,
access::address_space::global_space>(&pparent.send));
// Atomic increment allocates memory to write to.
lbegin[0] = cl::sycl::atomic_fetch_add(psstart_a, ltsum[0]);
// Atomic is necessary since multiple blocks access this
gbegin[0] = cl::sycl::atomic_fetch_sub(psend_a, gtsum[0]) - gtsum[0];

```

Note the creation of `cl::sycl::atomic<>` variables prior to the use of DPC++ atomic operations, which cannot operate on the global or local memory pointers directly.

So, let's summarize what we have so far: we translated and templated supporting structs and functions, converting specialized OpenCL C functions to DPC++ ones, created a template functor class with local accessors, and translated barriers and atomics. We move on to `lqsort_kernel` translation.

## Porting OpenCL Kernels to Data Parallel C++: Part 2 – `lqsort_kernel`

Translation of `lqsort_kernel` follows the now familiar patterns outlined by the translation of `gqsort_kernel`: we create `lqsort_kernel_class` functor to begin with and then translate local memory arrays and variables and barriers (it does not use atomics). `lqsort_kernel` also uses supporting function and structs. In addition to `plus_prescan` and `median_select` functions used by `gqsort_kernel`, we have `bitonic_sort` and `sort_threshold` functions that are considerably more complex and specific to `lqsort_kernel`. When translating them, they become the member functions of the `lqsort_kernel_class`. Their signatures change as well, since internally they use barriers, which in DPC++ case requires the use of iteration objects. They also work on local and global memory pointers, which require special handling, so the following OpenCL C signature:

```
void bitonic_sort(local uint* sh_data, const uint localid)
```

becomes

```
void bitonic_sort(local_ptr<T> sh_data, const uint localid, nd_item<1> id)
```

and

```
void sort_threshold(local uint* data_in, global uint* data_out,
    uint start, uint end,
    local uint* temp, uint localid)
```

becomes

```
void sort_threshold(local_ptr<T> data_in, global_ptr<T> data_out,
    uint start, uint end,
    local_ptr<T> temp_, uint localid,
    nd_item<1> id)
```

These functions are otherwise translated similarly to `gqsort_kernel`, with the `UINT_MAX` macro being replaced with `std::numeric_limits<T>::max()` to account for handling of various data types in the future.

When translating the `lqsort_kernel` itself, pointers to local memory, e.g. `local uint* sn;` are replaced with `local_ptr<>` objects, e.g. `local_ptr<T> sn;` Then to retrieve the local pointer from the local accessor, we call the `get_pointer` member function of the accessor object:

```
sn = mys.get_pointer();
```

`local_ptr<>` and `global_ptr<>` objects work with pointer arithmetic, so what previously was `d + d_offset`, where `d` was a global pointer, becomes

```
d.get_pointer() + d_offset
```

Note that local memory variables are translated as accessors of size 1, so they are translated as array accesses at index 0, e.g. `gtsum[0]`. When we are finished with the `lqsort_kernel` translation, we fully transitioned to DPC++, but we are still sorting unsigned integers. We did all the prework of templating supporting structs and functions and the functor classes of the two main kernel functions and now ready to enjoy the benefits.

## The Power of Data Parallel C++: Templates! And Their Caveats 😊

The real power of DPC++ is the ability to use C++ templates that are so near and dear to the heart of every C++ programmer. Templates enable programmers to write generic code. We want our GPU-Quicksort to be generic and to be able to sort not only unsigned integers but other basic data types as well, e.g. floats and doubles. We can, of course, make this sort even more generic and allow the user to overload comparison operator and handle composite types, but this remains an exercise for the reader. The main change in addition to `UINT_MAX` to `std::numeric_limits<T>::max()` change mentioned above is an additional modification of the `median_select` function. It turns out that `cl::sycl::select` takes a different type of the third argument depending on the size of the type of the first two arguments, so we need to introduce `select_type_selector` type traits class:

```
template <class T> struct select_type_selector;
```

```
template <> struct select_type_selector<uint>
{
    typedef uint data_t;
};
```

```
template <> struct select_type_selector<float>
{
    typedef uint data_t;
};
```

```
template <> struct select_type_selector<double>
{
    typedef ulong data_t;
};
```

It allows to convert a Boolean comparison to an appropriate type required by `cl::sycl::select`, so the `median_select` function becomes:

```
template <class T>
T median_select(T x1, T x2, T x3) {
    if (x1 < x2) {
        if (x2 < x3) {
            return x2;
        } else {
            return cl::sycl::select(x1, x3, typename select_type_selector<T>::data_t(x1 < x3));
        }
    } else { // x1 >= x2
        if (x1 < x3) {
            return x1;
        } else { // x1 >= x3
            return cl::sycl::select(x2, x3, typename select_type_selector<T>::data_t(x2 < x3));
        }
    }
}
```

If a user of the library wants to handle additional types, more specializations of `select_type_selector` should be added. Now, with the change above, our GPUQSort can sort floats and doubles on the GPU! Welcome to the DPC++!

## Back to Windows! And RHEL ...

The other powerful feature of DPC++ is its portability. To demonstrate it we will port the code to Windows and RHEL. The RHEL port is truly minimal: we need to add `imf` library at the link time. Windows port is only so slightly more complex: in addition to adding the following definitions when compiling:

```
-D_CRT_SECURE_NO_WARNINGS -D_MT=1
```

we also need account for the fact that `cl::sycl::select` for doubles requires `unsigned long long` type as the third parameter as opposed to `unsigned long` on Linux, so our `select_type_selector` for doubles becomes:

```
template <> struct select_type_selector<double>
{
#ifdef _MSC_VER
    typedef unsigned long long data_t;
#else
    typedef unsigned long data_t;
#endif
};
```

On Windows we need to undefine `max` and `min` to prevent the macro definitions from colliding with `std::min` and `std::max`. And that's pretty much it! So now we can sort unsigned integers, floats and doubles using Intel GPUs on Windows and two flavors of Linux with DPC++.

## Conclusion

It's been a long journey! We demonstrated how to gradually translate GPU-Quicksort from its original OpenCL 1.2 into DPC++. Note, that at every step along the way we had a working application, so when you are considering bringing DPC++ to your workflow, you can start small and slowly either add on or fully transition to DPC++ as your time allows. You can easily mix OpenCL and DPC++ in your code base and enjoy the benefits of both: you can still use legacy OpenCL kernels in their original form and enjoy the full power of C++ templates, classes and lambdas, when developing your new code in DPC++. You can easily port your code between Windows and various flavors of Linux and choose, which platform to develop on. You also have the full power of Intel tools with which to debug, profile and analyze your program. We hope that you enjoyed this article and will start using DPC++ in your projects!

## References

1. [Khronos OpenCL](#) The open standard for parallel programming of heterogeneous systems
2. [Khronos SYCL](#) C++ Single-source Heterogeneous Programming for OpenCL
3. [GPU-Quicksort: A practical Quicksort algorithm for graphics processors](#) by Daniel Cederman and Philippas Tsigas
4. [GPU-Quicksort in OpenCL 2.0: Nested Parallelism and Work-Group Scan Functions](#) by Robert Ioffe
5. [Intel® oneAPI Toolkits](#)
6. <Include the link to the code base here>